

Real-Time Automatic Gain Control for Singing Voice Applications

Bachelor Thesis in the course Informatik

Author:
Nils Heine

Matriculation Number: 6703759

Signal Processing / Signalverarbeitung
Department of Computer Science, MIN Faculty

First Reviewer: Prof. Dr.-Ing. Timo Gerkmann
Second Reviewer: Dr.-Ing. Martin Krawczyk-Becker

Hamburg, October 19, 2017

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Idea	1
2	Basics	3
2.1	Python	3
2.2	JUCE framework and C++	4
2.3	Test environment	4
2.4	Sources	4
3	Prototype	5
3.1	Overview	5
3.2	Filter	6
3.3	RMS	8
3.4	Time Coefficients	9
3.5	Gate	9
3.6	Gain	10
3.7	Lookahead	11
4	Optimisation	14
4.1	How	14
4.2	Circumstances	14
4.3	Approaches	15
5	Results	16
6	Discussion	17

7 Conclusion and Future Work	18
Bibliography	19
Index	19

List of Figures

3.1	Prototype processing chain	6
3.2	Plotted filter test	8
3.3	Gain adaption with and w/o lookahead	12
3.4	Initialized ring buffer with read and write pointers	13

List of Tables

CHAPTER 1

Introduction

1.1 Motivation

When a sound engineer edits a song he wants all the recorded audio tracks to be perceptible in the final mix (apart from some special cases). It is most important for audio tracks with notably significance for the musical piece. In this thesis I will work with vocal tracks due to their great significance in meaning, main melody or recognition value of the song. The difficulty with vocal tracks in the mix is the wide dynamic range that singers often use, unlike for example an distorted electric guitar which mainly stays on the same loudness level and is therefore easy to mix with great presence. Almost in every mix the vocals pass through an compressor to reduce their dynamic range. But this is rarely sufficient as compressors are working comparatively fast - too fast to compensate whole song parts or even some seconds with different vocal levels. For instance when a singer is changing his singing style or he sings instinctively quieter during an instrumental break which may not fit the mix. As a result it is a common procedure to automate an applied gain for every vocal track in the digital audio workstation (DAW) via sketching a gain curve by hand. Obviously this is a time consuming and monotonous task and therefore perfect to hand over to a machine.

1.2 Idea

The idea was to write a DAW plug-in that handles the former described problem. A plugin that will sketch a gain curve for a vocal track in real time. This will save the engineer time at every mixing session and the outcome will be more accurate as his own handwork.

As the recorded sound pressure level is not transferable to the perceived loudness, my goal was to adapt some algorithmic features of the ITU-R BS.1770-4¹ “algorithm to measure audio ... loudness”. This algorithm is an up-to-date standard for adjusting audio files to a same humanly perceived loudness level. This will help the gained audio track to stand out in the mix at every point.

¹International Telecommunication Union, Recommendation ITU-R BS.1770-4 (10/2015)

One plug-in with similar purpose called “Vocal Rider”² was published by Waves. My aim is to build my plug-in with features I consider important and useful for its purpose, while i do not know about the algorithm behind the Waves plug-in. Hence I will compare the resulting audio files after adapting the gain individually through the “Vocal Rider” and my own plug-in.

²[waves.com/plugins/vocal-rider](https://www.waves.com/plugins/vocal-rider)

CHAPTER 2

Basics

At the beginning of the development I testes ideas in python, later I used the JUCE¹ framework which is based on C++. The functionality of the plug-in was mainly tested in Logic Pro 9.

2.1 Python

I did not start with a final blueprint for the plug-in. Especially at the beginning I tried several ideas on the basic algorithm, the gate or loudness detection. In consequence I had to rearrange the code very often. So python came in handy as it focuses on code readability. In python code there are fewer steps necessary to write the same program as for example in C++ where the plug-in was finally written in.

Furthermore python provides various packages which extend its scope by useful features. For example the matplotlib.pyplot² plotting framework that ables you to draw graphs of your results. This was especially useful for testing on the filter implementation (see chpt. 3.2) and comparing optimisation results later on (see chpt. 4). The numpy³ package was essential for mathematical operations and the scipy⁴ tools very useful in terms of audio handling and optimization. For this thesis I used python version 3.6.

Still python was not my final choice for the plug-in as the C++ based JUCE framework offers a great predefined interface for audio plug-ins as well as the ability of fast processing due to the hardware-oriented C++ language. The Speed of calculations can be crucial for real-time audio processing. IRGENDWO NOCHMAL MEIN realtime erläutern (Introduction was es bedeutet!)

¹juce.com

²matplotlib.org/api/pyplot_api.html

³numpy.org

⁴scipy.org

2.2 JUCE framework and C++

JUCE is a cross-platform framework for audio applications based on C++. The main advantage for me is that it already contains the necessary functions for compiling to a working VST⁵ or AU⁶ plug-in. Therefore my main focus could stay on the algorithm of my plug-in during development. The JUCE audio plug-in template can be easily extended with a simple UI with sliders for the parameters of the algorithm. This is very useful for testing the effects of the individual parameters. JUCE takes over much of the communication with the DAW. Mostly this was fitting my plan and for this reason I just had to overwrite a few parts in which the plug-in had special needs.

2.3 Test environment

The JUCE framework brings along two ways to run your plug-in. The fastest one is to build the plug-in as standalone which can be done directly from the IDE⁷. The plug-in starts immediately and you can choose the main input and output channels. This is perfect for testing small bug fixes or visual changes. The standalone has its limitation in terms of for example a side chain input as it is not embedded in an surrounding DAW. But for this case JUCE has the Audio Plugin Host as solution. The Audio Plugin Host can host different plug-ins at the same time and visualises all inputs and outputs. It lets you draw connections between those ports and the currently active audio interface of the operating computer. The advantage over a real DAW is that you still get debugging output through runtime.

Still it is reasonably necessary to test it in a real DAW for a realistic environment and to be able to use all considered features for instance writing an automation or comfortably feeding a real backtrack into the side chain input. I used Logic Pro 9 to run my plug-in for the reason that it works with AU plug-ins which are per default supported by JUCE. Due to the custom UI it was still possible to change calculation parameters at runtime. Before there was a plug-in I have done basic algorithmic tests. Those I used python for, because of its simplicity and great visual possibilities. Therefore I could efficiently try different approaches and visualise if they have done their task correctly.

2.4 Sources

⁵Virtual Studio Technology plug-in architectur provided by Steinberg

⁶Audio Unit plug-in architectur provided by Apple

⁷integrated development environment

CHAPTER 3

Prototype

3.1 Overview

The basic approach works in five steps: filter, root mean square (RMS) calculation, gate, gain adaption, delay.

The processing chain starts with a low-cut and a high-shelf filter which will be applied sample wise on the incoming audio. These filters are a simplified mapping of the perception of sound pressure for human beings. This is the first step to adjust the plug-in to loudness instead of sound pressure (as done in FOOTNTE). However, the plug-in won't be running with the full loudness detection algorithm (FUTNOT) due to calculation time and real time capability (see SPÄTER? oder hier?).

After the filter section every sample is passed to the RMS calculation. The goal is to output the squared average of a previously specified period of time. The square root will be determined in the following calculation of the equivalent dB value. The plug-in converts the linear audio samples into the logarithmic dB scale in order to display gain values and loudness goal (see gain part?) in dB in the user interface (UI) as it is the standard scale of DAWs. Thereby the executive sound engineer will intuitively know how to interpret and interact with the UI (see design part?).

When the dB conversion is done, all the samples pass through an initially specified gate. The gate will set all samples with lower dB value as its threshold to the current loudness goal (see gain, see improve). In this way the plug-in will not operate when it is fed with silence or irrelevant noise.

Next step after the gate is the gain adaption. In this step the gated RMS value is compared to the current loudness goal. Depending on the difference of both values it will result in an preliminary gain. The gain variations per sample are smoothed comparable to the RMS calculation. This leads to the final gain value.

Lastly the new gain is multiplied with the current sample. Because the plugins behaviour is smoothed as it shall sound natural, it will not react instantly to the input. To compensate the reaction time it delays the input signal before multiplying the calculated gain. This delay is later offset by the DAW.

During development most of the parameters described in the following sections were settable in a basic dummy UI. This was realised with the standard JUCE slider and button objects and used for tests and fast adjustments.

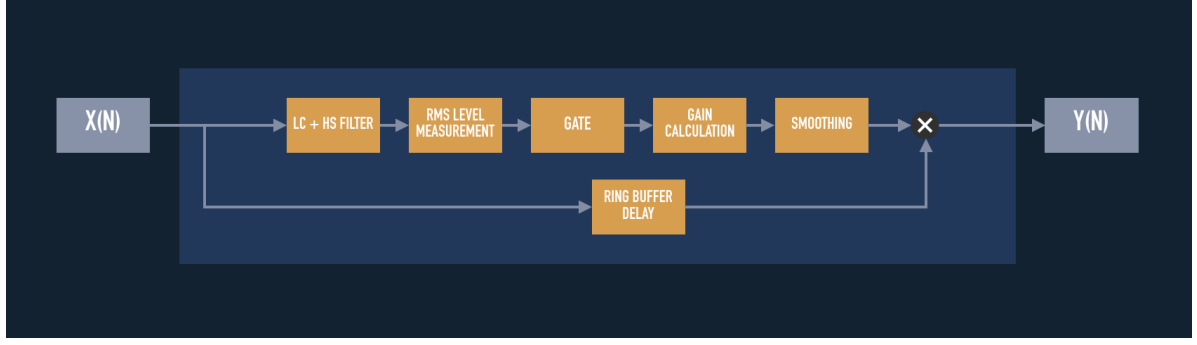


Figure 3.1: Prototype processing chain

3.2 Filter

While gathering information on how to improve my idea of the plug-in I got interested in the ITU-R BS.1770-4 [FOOTNOTE pls] algorithm. It classifies an audio file for its humanly perceived loudness. The main use of this algorithm is found in television and music streaming services as they can keep the program loudness steady while switching content. As the human perception is also of interest for mixing a song, I examined how it was done. Due to a good documentation about how to implement the loudness algorithm I build a copy of it in python and decided which elements I would adopt in my plug-in. The first elements were the two filters. Like described above, their use is to mimic the human perception of sound pressure at different frequencies. The first filter is a low cut for the reason that human hearing is insensitive to low frequencies. The second filter is a high shelf and is "used to account for the acoustic effects of the head" [wie zitiert man]. This imitation of the human hearing is greatly simplified but cost effective in terms of computation. The low cut filter has a cutoff frequency at 38 Hz, the high shelf around 1681 Hz. They are initialised at every plug-in startup in the JUCE method "prepareToPlay" with the current sample rate of the integrating DAW:

```
lowcut.setCoefficients(38.0, sampleRate, (1.0/2.0));
highshelf.setCoefficientsShelf(1681.0, sampleRate, 4.0);
```

The implementation is based on the biquad filter from the Book BLA (BLA NOTE auch im code). I have chosen the second order biquad filter architecture as it is a very flexible and simple solution with just two samples delay. The calculation of filter coefficients is adopted as follows:

f_c = cut-off frequency, f_s = sampling frequency (rate), $K = \tan(\pi f_c / f_s)$,
 Q = factor for height of the resonance, G = gain, $V_0 = 10^{G/20}$

Lowcut:

$$\begin{aligned} b_0 &= \frac{Q}{K^2Q + K + Q} & b_1 &= -\frac{2Q}{K^2Q + K + Q} & b_2 &= b_0 \\ a_1 &= \frac{2Q*(K^2 - 1)}{K^2Q + K + Q} & a_2 &= \frac{K^2Q - K + Q}{K^2Q + K + Q} \end{aligned}$$

Highshelf:

$$\begin{aligned} b_0 &= \frac{V_0 + \sqrt{2V_0}K + K^2}{1 + \sqrt{2}K + K^2} & b_1 &= -\frac{2(K^2 - V_0)}{1 + \sqrt{2}K + K^2} & b_2 &= \frac{V_0 - \sqrt{2V_0}K + K^2}{1 + \sqrt{2}K + K^2} \\ a_1 &= \frac{2(K^2 - 1)}{1 + \sqrt{2}K + K^2} & a_2 &= \frac{1 - \sqrt{2}K + K^2}{1 + \sqrt{2}K + K^2} \end{aligned}$$

Hence the loudness algorithm uses second order filters (with two delay memories) it works like this:

Biquadfilter Bild

It is the same as my implementation in the Filter class:

```
double AutoVocalCtrlFilter::process(double sample)
{
    const double mid = sample - a1 * z_1 - a2 * z_2;
    const double out = b0 * mid + b1 * z_1 + b2 * z_2;

    z_2 = z_1;
    z_1 = mid;

    return out;
}
```

Before implementing in C++ I testen my filter class in python. Therefor I send different signals with frequencies between 0 and 20000Hz through both filters and plotted the resulting amplitudes in a graph via pyplot(alter fußnoten verweis?). The current algorithm results in a descent graph (Fig. 3.2 NOCH ÄNDERN). To test the C++ version of the filter I compared the results of the same input with the previously tested python implementation.

The implementation is capable of many filter styles at different cutoff frequencies. For my plug-in it is used for the low cut and high shelf filter described above, which are simply processed one after the other on the current audio sample for each channel Wie VIELE AM ENDE??:

```
double updateFilterSample(double sample, AutoVocalCtrlFilter hs,
AutoVocalCtrlFilter lc)
{
    return hs.process(lc.process(sample));
}
```

The updateFilterSample method returned the filtered samples for further processing.




Figure 3.2: Plotted filter test

3.3 RMS

To measure the level of the audio signal at the current position the plug-in calculates a root mean square (RMS) value of the last XX ms. This means it is calculating the root of the average from the squares of each sample in the set time window.

It uses RMS calculation because it is part of the imitation of human perception and because it has the necessary time to calculate an average value as it does not need to react fast. A human will (F. A. Everest, *The Master Handbook of Acoustics*, New York: McGraw-Hill, 2001.) not interpret a small impulse of a handful of samples as loud as a audio signal at the same level of longer duration. For example “a 3-msec pulse must have a level about 15 dB higher to sound as loud as a 0.5-sec (500-msec) pulse”. This is also part of the BLABLA loudness standard. Additionally the “RMS is equal to the value of the direct current that would produce the same average power in a resistive load” [wikipedia].

The RMS implementation is based on the Book *Digital Audio Signal Processing* by Udo Zölzer (123) and can be performed in one line of code:

```
double updateRMS2(double sample , double last , double co)
{
    return (1. - co) * last + co * (sample * sample);
}
```

The already filtered samples are used for this calculation and the result is written into

the rms2 vector. The Co parameter is a time coefficient (see gleich) of small size. This parameter influences the weight on how much the current squared sample will inflict the quadratic mean. In normal RMS procedure the final result is the square root of the value the plugin results in. In this case we can skip the calculation because it will be converted into the logarithmic dB scale in the next step and a square just changes the value of one coefficient in this process. The final RMS value will only appear as logarithmic dB but save a little time.

3.4 Time Coefficients

The time coefficients are used for RMS calculation and to realise the compress and expand times (see later). They are determined with a formula by Udo Zölzer (quell):

$$\begin{aligned} 1.f - \exp(-2.2 * (1./currentSampleRate)/(ms/1000.)) \\ = 1.f - \exp(-2.2 * (1./currentSampleRate)/s) \end{aligned}$$

Udo Zölzer decided to use the exponential function because it draws a natural decay. He determined -2.2 for the first part in the exponent by solving an equation system to achieve an attack time MATHE $t_a = t_{90} - t_{10}$. This means that when a calculation similar to the RMS calculation defined above uses such an time coefficient, its reaction on a input change will need the chosen amount of time to get from 10The second part in the exponent $(1./currentSampleRate)/s$ calculates the proportion of one sample to the amount of seconds of the MATH t_a . This needs to be done because it will be used on every sample.

BILD MIT ATTACK ZEIT ODER RMS MITTEL DURCH IMPLUSE VERÄNDERUNG
 t_{90} und t_{10}

3.5 Gate

The plug-in should operate while there are vocals and stop if there are none or just a soft decay. Else it would produce unwanted effects by amplifying noise. Additionally it would distort the applied gain for the actual vocals through strong gain increase at the gaps in-between.

To solve this potential problem the plug-in uses a gate. The gate checks for every sample if the sound pressure level is over a certain threshold. If not, it will be replaced by the current loudnessGoal. As the threshold and the loudnessGoal are set in dB, the first step in the gate is to convert the transferred sample (rms^2) into the logarithmic scale.

It uses $10.0 * std :: \log_{10}(rms^2 + 1e - 10)$ instead of $20.0 * std :: \log_{10}(rms + 1e - 10)$ because the rms2 is squared. It adds $MATH + 1e - 10$ to head off the undefined $\log_{10}(0)$ case. After the conversion it gates the dB rms sample value at the dB threshold.

The threshold is defined as `loudnessGoal - gainRange` IST SO GEBLIEBEN?. Thereby it is still possible to use the whole `gainRange` for gain adaption and at the same time sort out the decay of the vocals. With this formula the threshold adjusts to the level of the vocals as the `loudnessGoal` is detected (see später).

When rms samples are not passing the gate and therefore being replaced by the `loudnessGoal` the plug-in adapts the gain to 0 (after a short period of time) because it has achieved its goal (see `gain`).

3.6 Gain

Now the most crucial part is happening: the calculation of the final gain value for the current sample.

```
double updateGain(double sample , double lastGn)
{
    const double g = *loudnessGoal - sample;
    const double co = g < lastGn ? compressTCo : expandTCo;
    updateAutomation();
    return clipRange.clipValue((1 - co) * lastGn + co * g);
}
```

At first it computes the difference between the `loudnessGoal` and the current processed sample. The result is the gain factor that would be necessary to get it to the `loudnessGoal` (by multiplying in the linear number space).

Since the plug-in is designed to react on loudness differences for longer duration than for example compressors or expanders do, the gain adaption is smoothed over a proportionate amount of time. The smoothing is attained similar to the RMS calculation but uses different time coefficients.

The fitting time coefficient is chosen by comparing the calculated gain factor `g` with the gain that the function had returned for the precious sample. If `g` is smaller than the last gain (`lastGn`) the gain for the current sample will be smaller too. Therefore the dynamics of the input vocals will be compressed in relation to the last processed sample so it chooses the `compressTCo` time coefficient. The other way round when dynamics are expanded the `expandTCo` time coefficient is chosen. Different time coefficients for expanding and compressing dynamics are useful as amplifying a signal can be risky, while compressing it causes no trouble. For instance boosting a weak signal also boosts all the recorded unwanted noise. Additionally digital dynamics are limited and as the signal expands it risks to clip at the 0dB cap and produce distortion.

EXPAND VIELLEICGHT SCHWEIRIGES WORT DAFÜR WEIL SCHON BESETZT
vielleicht leise geht unter lautes sticht auffällig raus

After calculating the smoothed gain it gets clipped at the user chosen range up to $\pm 10\text{dB}$. On one hand this ables the user to adjust the maximum variation of dynamics on the other hand it prevents the gain to increase up to problematically high values. This does not happen during normal use in an expected environment but can't be completely ruled out due to possible unknown software errors or an unknown environment. As an error in the adapted gain does not only affect further calculation results but also the mixing engineer who is listening to an amplified signal, it is of special importance to avert wrong values (see LAST CAP OF PLUGIN).

When the gain is finally determined it will be converted back to the linear number space. Now it just needs to be multiplied with the current sample.

3.7 Lookahead

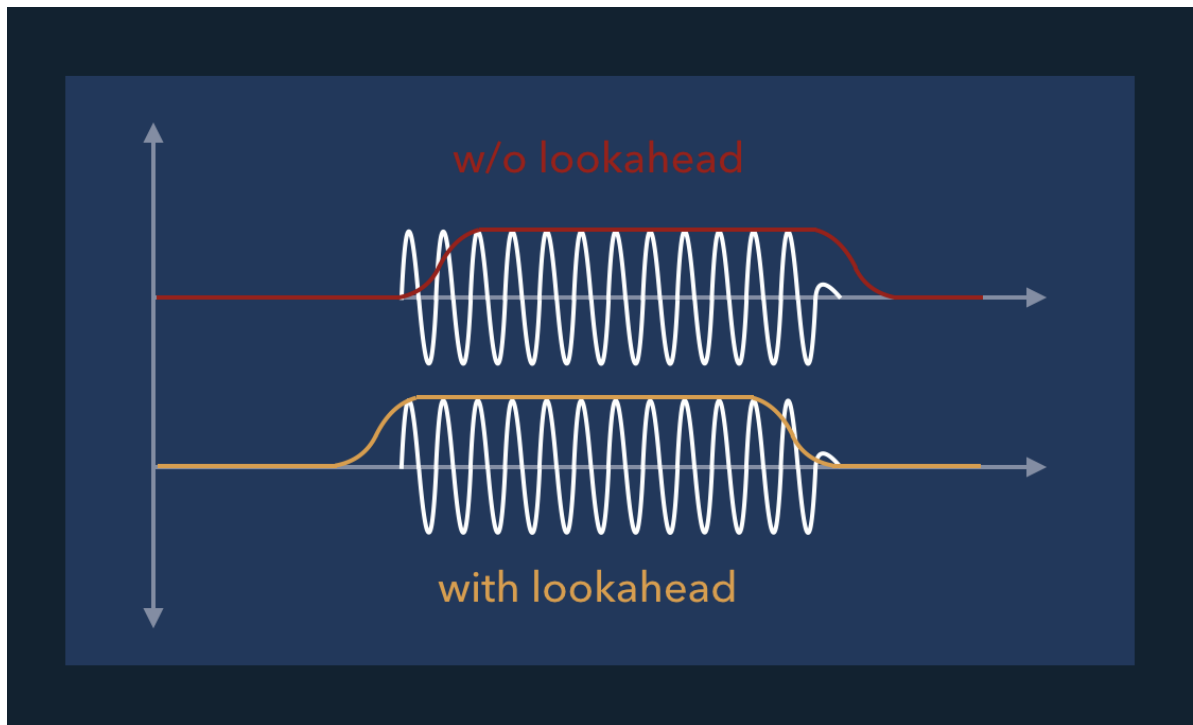


Figure 3.3: Gain adaption with and w/o lookahead

To achieve its goal the plug-in is designed to react slow and therefore needs some time adapt on a change in the average signal level. If it multiplies every sample with the gain value calculated form the same sample, the plug-in will be adjusted after the RMS has

changed to the new average and the smoothed gain has slowly adapted. This will take some ms and consequently the first part of the alternate signal is not perfectly gained. While in particular this part introduces a new section in a song for example a chorus, it is desirable to have an adapted gain already at this point.

To compensate the adaption time I implemented a simple lookahead feature which allows the plug-in to calculate the gain for the current samples while looking at future samples. This is realised by a ring buffer with two pointers at different locations. One write pointer to write the current sample transferred from the DAW into the buffer which is also used to determine the gain adaption and one read buffer ahead of it which is pointing on the sample that will be multiplied with the determined gain. To make this possible the gap between both pointers is as large as the set samples of the lookahead (converted from ms) and filled with zeros at the initialisation of the plug-in. In order to avoid that the whole plug-ins output is delayed I use the `setLatencySamples()` (see code besipel) method from JUCE to communicate the resulting delay with the embedding DAW. Therefore a correct woking DAW will send the signal earlier to the plug-in and the output remains at the correct position despite the lookahead.

For the development I realised a fader in the UI to adjust the lookahead at runtime but in the final build the amount of samples is constant as I want the plug-in to be as simple to work with as possible. naja und weil ich ja gewählt habe aus grund:

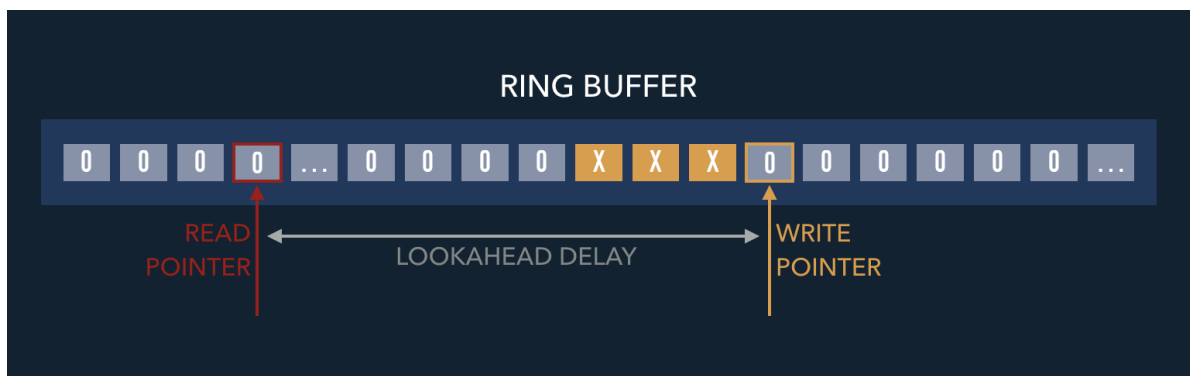


Figure 3.4: Initialized ring buffer with read and write pointers

vocals wichtig am Anfang erster push sollte nicht leiser gemacht sein weil vorher laut war oder andersrum finale größe von delay sagen LOL sollte nicht größer als ums sein oder verweise zu den Bildern und code

```
...
gain[channel] = updateGain(updateGate(rms2[channel], newGate),
gain[channel]);
double g = pow(10, gain[channel]/20);
delayData[dpw] = channelData[sample];
...
const double o = delayData[dpr] * g;
```

```

...
channelData[sample] = o;
...
if (++dpr >= delayBufferLength)
    dpr = 0;
if (++dpw >= delayBufferLength)
    dpw = 0;
...

```

```

void AutoVocalCtrlAudioProcessor::updateDelay()
{
    int delayInSamples = msToSamples(*delayLength);
    delayReadPos = (int)(delayWritePos - delayInSamples
+ delayBufferLength) % delayBufferLength;
    setLatencySamples(delayInSamples);
}

```

CHAPTER 4

Optimisation

After finishing the prototype of the plug-in I was interested in comparing the main functionality with the equivalent from WAVES. I wanted to know how likely the results can get with fitting parameters at my version of the gain adaption. Not only to see if I may have forgotten a important part so far but also to find out how flexible my plug-in is and to gather thoughts about how to set my own parameters and constants later on.

4.1 How

It would not be very effective to try different picks for the parameters by hand and compare the outcome as there are at least a handful of parameters which can be adjusted to a huge amount of possible combinations. Therefore I had to give this task to the computer.

Conveniently there is the `scipy.optimize` package for python which deals with optimisation tasks for example the algorithmic minimisation of a problem according to the result of a self defined function.

To make use of this package I primarily transferred the current code of the plug-in from C++ back to python were I just tested algorithmic ideas so far. After the python duplicate was ready to use I wrote a function to compare the resulting audio files after processed with the “Vocal Rider” and my own plug-in. It returns a value describing the deviation.

Through the optimisation process the parameter adjustments of the “Vocal Rider” differed between attempts but stayed constant in the process. The parameters of my plug-in were changed continuously to achieve a preferably small deviation.

4.2 Circumstances

In order to let the optimisation algorithm have the option of adjustment on different parts of my plug-in, I declared the `loudnessGoal`, `RMS time`, `compress time`, `expand time`, `gate` and `lookahead` as variables. At start I set a guessed values for each of the

parameters and set them in a array which was altered thru optimisation process and fed into my deviation function at every step of it.

For measuring the deviation for the current parameter array my function sums up the squared difference between both resulting audio files at every sample. The result from the “Vocal Rider” was therefore created in the DAW Pro Tools 11. My own implementation is called in the deviation function with the current parameter array.

When the optimisation search is done I feed the outcome into another function which displays the gain adaption from each plug-in in a collective graph. This gives a great overview of the result and remaining diversion of the implementations.

4.3 Approaches

The first tries unfortunately did not achieve a reasonable solution. Algorithmically the optimisation tests small variations in the parameter array and watches the outcome. If the initial guess is not well set, the differences are of very little amount in my comparison function and the `scipy.minimize` algorithm will not know were continue its search. WAUM=?

CHAPTER 5

Results

CHAPTER 6

Discussion

In the gggg

CHAPTER 7

Conclusion and Future Work

In this thesis, we yadaaaaaa

Ich stimme der Einstellung der Arbeit in die Bibliothek des Fachbereichs Informatik zu.

Ort, Datum

Unterschrift

Hiermit versichere ich an Eides statt, dass ich die vorliegende Arbeit im Masterstudengang Informatik selbstständig verfasst und keine anderen als die angegebenen Hilfsmittel — insbesondere keine im Quellenverzeichnis nicht benannten Internet-Quellen — benutzt habe. Alle Stellen, die wörtlich oder sinngemäß aus Veröffentlichungen entnommen wurden, sind als solche kenntlich gemacht. Ich versichere weiterhin, dass ich die Arbeit vorher nicht in einem anderen Prüfungsverfahren eingereicht habe und die eingereichte schriftliche Fassung der auf dem elektronischen Speichermedium entspricht.

Ort, Datum

Unterschrift